

WebAssembly GC Optimization Update

Alon Zakai

Aug 23, 2021





Big Picture

Good benchmark results on Dart.

Java (J2CL) results **not** good enough yet.


Wasm GC being fast is a blocker for shipping - or else, why wouldn't people use JavaScript?



Talk Focus

Mostly on the toolchain.

Need to do the things that both Closure Compiler and the JS VM were doing - much can only be done in the toolchain.





Main Theme

Not really “nominal vs structural.” The key issue is: “**more static type info vs less.**”

Wasm already has lots of useful static type info (e.g. type info for function parameters and results) and is getting even more (e.g. per-table types).



Devirtualization

All calls in Java are **indirect** (virtual). Indirect calls are **not fast**.

Try to turn them into **direct** calls, which can then also be **inlined** & further optimized.



```
(type $A (struct (ref $A.vtable) i32))  
(type $B (struct (ref $B.vtable) i32))  
(type $C (struct (ref $C.vtable) i32 f64)  
  (extends $A))
```

```
(type $A.vtable (struct (ref $functype.A)))  
(type $B.vtable (struct (ref $functype.B)))  
(type $C.vtable (struct (ref $functype.C)  
  (ref ..))  
  (extends $A.vtable)))
```

- Differentiating **\$A** and **\$B** (& their vtables) is useful.
- Knowing that **\$C** extends only **\$A** is useful.



Devirtualization: Impl

- In the entire module, find all the values written to struct fields, like
`(struct.new $A.vtable (ref.func $foo))`
or in `struct.set`, hoping for **constants**.
- Find `struct.gets` and figure out which `struct.new` and `struct.sets` are **relevant** to it.
 - If only a single constant value, **apply it!**



Static Type Info Helps!

```
(type $A.vtable (struct (ref $functype.A)))  
(type $B.vtable (struct (ref $functype.B)))  
(type $C.vtable (struct (ref $functype.C))  
                  (extends $A.vtable)))
```

```
;; affects $A.vtable, $C.vtable, not $B.vtable  
(struct.set $A.vtable ..)
```

```
;; affects $B.vtable, not $A.vtable, $C.vtable  
(struct.set $B.vtable ..)
```




Devirtualization: Results

Using `--nominal` in Binaryen:

- **41%** speedup.
- **4%** code size reduction.



Cast Elimination?

Some casts needed for validation. Others, not (imagine this is after inlining or such):

```
(func $foo (param $x (ref $A))  
           (result   (ref $A))  
  (ref.cast  
    (local.get $x)  
    (global.get $A.rtt)))
```



Cast Elimination?

Even if we write `rtt.canon` there, we can't remove the cast.

```
(func $foo (param $x (ref $A))  
           (result (ref $A))  
    (ref.cast  
      (local.get $x)  
      (rtt.canon $A))) ;; changed
```



Cast Elimination?

Binaryen introduced a new **TrapsNeverHappen** mode. If set, the optimizer assumes the program will not actually trap.

This is a **13%** speedup and **6%** code size reduction.

```
(func $foo (param $x (ref $A))  
           (result    (ref $A))  
    (local.get $x)) ;; no cast!
```



Cast Elimination?

Even `TrapsNeverHappen` does not help with `ref.test`. But **static type information** would!


```
(func $foo (param $x (ref $A))  
           (result (ref $A))  
    (ref.test $A ;; => (i32.const 1)  
      (local.get $x)))
```



Toolchain-Only?

Wasm object files today have more information than normal object files, like **relocations**.

We could in principle do the same in wasm GC:





Reasons Against Toolchain-Only

Relocations don't make sense for shipping wasm. But static GC type info does: **Inlining in the VM** for GC, for example.

Simpler overall ecosystem.





Conclusion

More static type information is good!

This is **not** in opposition to structural typing, so long as structural behavior is **in addition**.

Questions?