



# WebAssembly Exception Handling (Phase 1)

Heejin Ahn and Michael Starzinger

# Involved parties

- Champion / WebAssembly Toolchain Implementation: Heejin Ahn
- Google WebAssembly Toolchain TLM: Derek Schuff
- Google WebAssembly V8 TL / V8 Implementation: Michael Starzinger

# Agenda


- Proposed goals
- Spec recap
- Status updates
- Discussion: Should we catch traps with `catch`?

# Proposed goals

- Provide a primitive for exception handling for Wasm programs
  - Zero-cost: no runtime cost until exceptions are thrown or caught
  - Structured: composes properly with existing control flow constructs
  - Safe: fast, single-pass verification
- Low-cost C++ exception handling
  - Emscripten currently indirections through JS for try/catch (expensive, not production quality)
  - WASM solution should be nearly as efficient as a native target

# Event section

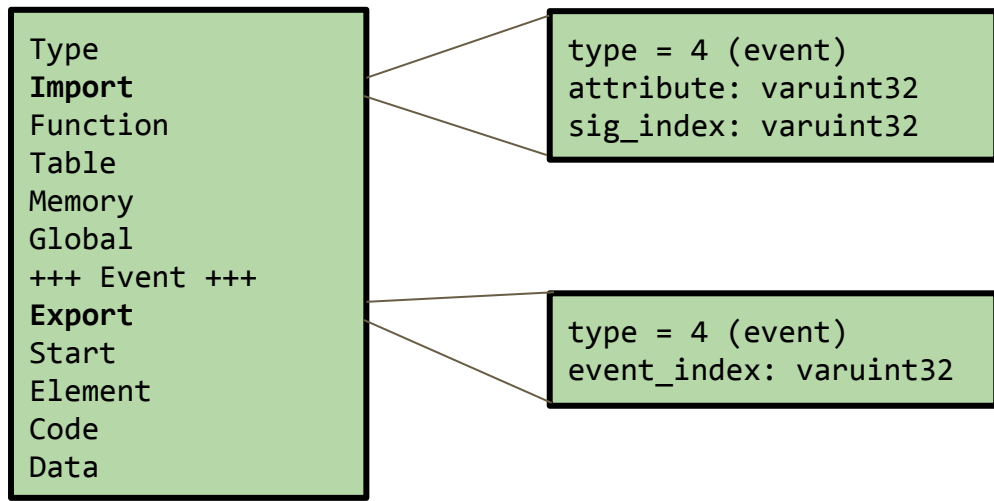
- Wasm events are features that suspend the current execution and transfer the control flow to a corresponding handle
  - Only supported kind is exceptions now



- Event section declares a list of event types
- `attribute`: application-specified number (0 for exceptions)
- `sig_index`: index into types section of function signature

# Event import/export


- Events can be imported and exported from a module
- Each instantiation of a module produces new event tags



- Imported events specify an attribute and an expected signature
- Exported events specify an exported event index

# Throwing an exception

- A **throw** accepts its arguments on the stack and creates an **except\_ref** value out of them
- and begins searching the *control* and *call* stacks for a handler



# The exception reference data type

- EH proposal requires the [reference types proposal](#) as a prerequisite
- `except_ref` type is a subtype of `anyref`
- `except_ref` type contains values thrown
- `except_ref` type possibly can contain more information, such as a stack trace



# try and catch blocks

- **try ... catch ... end** introduces a new block kind
  - **try** can have a label for branches too
- If any instruction between a **try** and a **catch** throws, the VM unwinds the wasm execution stack and resumes execution from the **catch**
- A **catch** pushes an **except\_ref** value onto the stack

```
try [block_type]
  ...
catch
  ...
end
```

# Rethrowing an exception

- A **rethrow** takes an **except\_ref** value from the stack (produced by a **catch**) and continue unwinding the execution stack with the exception
- A **rethrow** can occur anywhere (not necessarily between **catch** and **end**)

# Exception data extraction

- A `br_on_exn` checks the exception tag of an `except_ref` on top of the stack (without popping it) if it matches the given exception index
  - If they match, it
    - branches out to the label referenced
    - pops the `except_ref` from the stack
    - extracts the `except_ref` and pushes the exception's values onto the stack
  - If they don't match, it does nothing, and the `except_ref` remains on the stack
- Format: `br_on_exn label except_index`

# Exception data extraction

- We can also test an `except_ref` against multiple tags

```
block $l0 (result i32, i64)
  block $l1 (result i32)
    ...
    ;; except_ref $e is on the stack at this point
    br_on_exn $l1 e(i32)          ;; branch to $l1 with $e's arguments
    ;; except_ref $e is left on the stack if br_on_exn is not taken
    br_on_exn $l0 e(i32, i64) ;; branch to $l0 with $e's arguments
    rethrow
  end
  ;; handler for $l1
end
;; handler for $l0
```

# Toolchain implementation status

- Current status

- LLVM part is mostly done, modulo proper `throw`; keyword support
- Other toolchain support in progress (Binaryen / emscripten)

- Next steps: short term

- Finish toolchain and libcxxabi support
- Start end-to-end testing

- Next steps: long term

- Support `throw`; keyword properly
  - Requires full reference type proposal support, including exception tables to maintain a stack of in-flight `except_refs`
- `setjmp-longjmp` handling
- Benchmark tests and optimizations

# V8 Implementation Status (1)

- Current status
  - Full support of current proposal in top-tier and interpreter (for debugging)
  - Behind the `--experimental-wasm-eh` runtime flag
- Interaction with other language features
  - Support for `except_ref` in signatures, locals, and globals, and (partially) tables.
- Implementation details
  - `try ... catch ... end` implemented as “zero-cost” for non-exceptional execution
  - Exception packages (`except_ref` values) allocated on garbage collected heap (`throw` allocates)

# V8 Implementation Status (2)

- Interaction with JavaScript embedding
  - Exported exceptions/events are an instance of `WebAssembly.Exception`
  - Exception packages are instances of `WebAssembly.RuntimeError`
  - Exception package contain full stacktrace (WebAssembly & JavaScript frames)
  - No access to exception tags/values encoded in package from JavaScript
  - Cannot construct WebAssembly exceptions from JavaScript (neither package nor event)
  - All JavaScript exceptions caught by `catch`, can be `rethrown` but never match `br_on_exn`
  - In JavaScript `anyref` and `except_ref` are indistinguishable (can be any JavaScript value)
- Current interaction with traps (self-consistent but undesirable)
  - `try ... catch ... end` will catch traps only if it crosses one or more function invocations
  - Conceptually: traps are converted into exceptions at function boundary

# Should a catch catch traps?

```
try
  local.get 0
  local.get 1
  i32.div_s      ;; possibly divide by zero
  ...
catch
  ;; recover from the trap??
end
```



# What kind of traps do we have?

- Undefined arithmetics (e.g. divide by zero)
- Invalid memory/table accesses
- Ill-typed indirect function calls
- GC-related traps

# Reasons for catching traps?

- Error recovery and continuing execution / Backtrace preservation
  - Which traps can be recovered from and how?
- Invariant: All non-local control-flow transfers can be intercepted by catch
  - Otherwise traps would represent a non-local control-flow transfer from WebAssembly to the embedder (e.g. JavaScript) that is impossible to notice/intercept by a surrounding function when functions/modules are composed

# Caveats

- Catching traps in wasm does NOT mean languages' **catch** clauses do too. Then how useful would catching traps be?
  - Some languages' (e.g. C++) **catch** clauses don't catch traps
    - In case someone wants to do that, it's hard
  - Some languages (e.g. Rust or Java) turn trap-like behaviors to exceptions, so wouldn't be useful for those
    - Compiler-generated code include checks for trappable instructions
- We need extra checking code after every **catch** instruction
  - *"Is this a trap? Then don't run destructors / cleanup code and just rethrow"* per every catch block, including all landing pads created by destructor frames

# Representation of traps

- Currently in browsers wasm traps surface as JS exceptions, from which point they become 'foreign exceptions' which are caught
  - Whether we catch traps or not, we should at least be consistent before we reach embedder and after
- We need a way to distinguish traps from foreign exceptions anyway
  - Special trap tag?



# Discussions

# Backup Slides

# Want to make C++'s catch clause catch traps?

- Every throwable call ends a BB and has a normal exit and an unwind exit, but other trappable instructions don't
- After CFG transformation, we don't guarantee traps end up in the correct catch clause

```
try {  
    foo();  
    bar();  
} catch (...) {  
}
```



# Want to make C++'s `catch` clause catch traps?

1. High-overhead solutions: Not zero-cost anymore
  - Solution a: Insert unwind edges after all other trappable instructions
  - Solution b: Outline every region between a `try` and a `catch` clause
    - Windows SEH takes this approach
2. Compiler inserts code so that after `catch` instruction, if it's a trap, always rethrow it until we reach the top frame
  - Involves small code size overhead of checking and rethrowing after every `catch` instruction
  - How is it different from not catching traps?