

WASM Trace Proposal

July 20, 2021

Jacob Abraham and Rich Winterton

Proposal

- “instTrace” custom section
- Contains a list of trace instructions
 - Each instruction is encoded as a byte offset into the code section followed by an immediate ID
 - Immediate ID allows traces to be identified easily
 - No need to specify instruction in WASM
 - Requires the target implementor to know which native code sequence
- Intrinsic
 - `__builtin_wasm_trace_instruction(ID)`
 - Prototype in clang
- Runtime
 - Decode and inserted in instruction stream
 - Prototype in V8

C Source Example

- Should add a custom section
 - 2 traces, one before the i32.add and one after

```
int add(int a, int b) {  
    __builtin_wasm_trace_instruction(17);  
    int c = a+b;  
    __builtin_wasm_trace_instruction(18);  
    return c;  
}
```

ASM Example

```
clang --target=wasm32 -c test.c -S -O3 -o test.s
```

What the prototype generates

```
add:
    .functype    add (i32, i32) -> (i32)
.Ltrace0:
    local.get    1
    local.get    0
    i32.add
    local.set    1
.Ltrace1:
    local.get    1
    end_function
```

Custom section defining the labels

```
.section .custom_section.instTrace,"",@
.int8    2
.int32    .Ltrace0
.int8    17
.int32    .Ltrace1
.int8    18
```

WASM Example

```
clang --target=wasm32 -c -emit-llvm test.c -O3 -o test.bc
```

```
llc -march=wasm32 -filetype=obj test.bc
```

```
wasm-ld --no-entry --export=add test.o -o test.wasm
```

Ideally what we want

Contents of section Custom:

0969 6e73 7454 7261 6365	len, instTrace
0203 110a 12	num traces,
	offset, immediate,
	offset, immediate

Ideally, we want uvar32 for
offsets, limitation in clang

What the prototype generates

Contents of section Custom:

0969 6e73 7454 7261 6365	len, instTrace
0203 0000 0011 0a00 0000 12	num traces,
	offset, immediate,
	offset, immediate

*discussed in issues

Runtime Code Gen

Experiment running in d8, using only Liftoff

TurboFan work in progress*

d8 --experimental-wasm-instruction-tracing --print-code --liftoff-only test.js

With runtime flag “--wasm-trace-native=<type>”, we can switch what native code is run for different simulators.



Minimal overhead, on par with optimized debug builds


*discussed in issues

00000350B49413C0	0	55	push rbp
00000350B49413C1	1	4889e5	REX.W movq rbp, rsp
00000350B49413C4	4	6a08	push 0x8
00000350B49413C6	6	4881ec08000000	REX.W subq rsp, 0x8
00000350B49413CD	d	488975f0	REX.W movq [rbp-0x10], rsi
00000350B49413D1	11	488b5e23	REX.W movq rbx, [rsi+0x23]
00000350B49413D5	15	483b23	REX.W cmpq rsp, [rbx]
00000350B49413D8	18	0f864a000000	jna 00000350B4941428 <+0x68>
00000350B49413DE	1e	53	push rbx
00000350B49413DF	1f	bb11000000	movl rbx, 0000000000000011
00000350B49413E4	24	6467909090	sscmack
00000350B49413E9	29	5b	pop rbx
00000350B49413EA	2a	8d1c02	leal rbx, [rdx+rax*1]
00000350B49413ED	2d	53	push rbx
00000350B49413EE	2e	bb12000000	movl rbx, 0000000000000012
00000350B49413F3	33	6467909090	sscmack
00000350B49413F8	38	5b	pop rbx
00000350B49413F9	39	8bc3	movl rax, rbx
00000350B49413FB	3b	48837df808	REX.W cmpq [rbp-0x8], 0x8
00000350B4941400	40	7421	jz 00000350B4941423 <+0x63>
00000350B4941402	42	b924000000	movl rcx, 0000000000000024
00000350B4941407	47	4989e2	REX.W movq r10, rsp
00000350B494140A	4a	4883ec28	REX.W subq rsp, 0x28
00000350B494140E	4e	4883e4f0	REX.W andq rsp, 0xf0
00000350B4941412	52	4c89542420	REX.W movq [rsp+0x20], r10
00000350B4941417	57	48b84043eb25f67f0000	REX.W movq rax, 00007FF625EB4340
00000350B4941421	61	ffd0	call rax
00000350B4941423	63	488be5	REX.W movq rsp, rbp
00000350B4941426	66	5d	pop rbp
00000350B4941427	67	c3	retl
00000350B4941428	68	50	push rax
00000350B4941429	69	51	push rcx
00000350B494142A	6a	52	push rdx
00000350B494142B	6b	56	push rsi
00000350B494142C	6c	e84ffefffff	call 00000350B4941280
00000350B4941431	71	5e	pop rsi
00000350B4941432	72	5a	pop rdx
00000350B4941433	73	59	pop rcx
00000350B4941434	74	58	pop rax
00000350B4941435	75	eba7	jmp 00000350B49413DE <+0x1e>
00000350B4941437	77	90	nop

DevTools Enabling

- SDE (Software Development Emulator)
- Browser enabling
 - Debug breakpoints
 - Start/Stop performance trace

	sde-mix-out.txt	14,016 KB
	sde-mix-out-partial.txt	103 KB

 Debug?	Mark	Location
<input type="checkbox"/>	17	Line 3 in ee9051d6
<input type="checkbox"/>	18	Line 3 in 18b8d8c6
<input type="checkbox"/>	18	Line 4 in 18b8d8c6
<input checked="" type="checkbox"/>	2	Line 3 in b17e684a
<input type="checkbox"/>	2	Line 8 in b17e684a

b17e684a x	
0x000	(module
0x01d	(func \$any (;0;) (export "any")
0x020	ittnop 2
0x023	i32.const 10
0x025	i32.const 7
0x027	i32.add
0x028	drop
0x029	ittnop 2
0x02c)
0x02d)

Current Issues With Prototype

- V8
 - In Liftoff, trace instruction is emitted where the offset is.
 - In TurboFan, inside of a function if there are more than one trace instruction, all but the first get optimized away.
- Emscripten
 - The compilation works fine (just clang), but the optimization step with binaryen sometimes causes problems. When the WASM gets reordered, the instTrace offsets don't move with it.

Summary

- Instruction Tracing
 - Custom section defines where trace instructions should be emitted
 - Decision made at the engine implementation
 - Compiler intrinsic to generate custom section
- <https://github.com/WebAssembly/instrument-tracing>
- Vote for Phase 2
- Next Steps in Phase 2
 - Support from engine teams to integrate prototype
 - Support from tools teams
 - Clang uvar32 support
 - Support for development on DevTools
 - Feedback from early adopters, including additional testing