

WASM Trace Proposal

July 20, 2021

Jacob Abraham and Rich Winterton

Proposal

- “instTrace” custom section
- Contains a list of trace instructions
 - Each instruction is encoded as a byte offset into the code section followed by an immediate ID
 - Immediate ID allows traces to be identified easily
 - Don’t need to specify instruction in WASM
 - Requires the target implementor to know which native code sequence
- Intrinsic
 - `__builtin_wasm_trace_instruction(ID)`
 - Prototype in clang
- Runtime
 - Decode and inserted in instruction stream
 - Prototype in V8

C Source

- Should add a custom section
 - 2 traces, one before the i32.add and one after

```
int add(int a, int b) {  
    __builtin_wasm_trace_instruction(17);  
    int c = a+b;  
    __builtin_wasm_trace_instruction(18);  
    return c;  
}
```

ASM

```
clang --target=wasm32 -c test.c -S -O3 -o test.s
```

Ideally what we want

```
add:
    .functype    add (i32, i32) -> (i32)
.Ltrace0:
    local.get    1
    local.get    0
    i32.add
.Ltrace1:
    end_function
```

What the prototype generates

```
add:
    .functype    add (i32, i32) -> (i32)
.Ltrace0:
.Ltrace1:
    local.get    1
    local.get    0
    i32.add
    end_function
```

*discussed in issues

Custom section defining the labels

```
.section .custom_section.instTrace,"",@
.int8    2
.int32   .Ltrace0
.int8    17
.int32   .Ltrace1
.int8    18
```

WASM

```
clang --target=wasm32 -c -emit-llvm test.c -O3 -o test.bc
```

```
llc -march=wasm32 -filetype=obj test.bc
```

```
wasm-ld --no-entry --export=add test.o -o test.wasm
```

Ideally what we want

Contents of section Custom:

```
0969 6e73 7454 7261 6365      len, instTrace
0205 0000 0011 0a00 0000 12  num traces,
                                offset, immediate,
                                offset, immediate
```

What the prototype generates

Contents of section Custom:

[illegible]

Runtime Code Gen

Experiment running in d8, using only Liftoff

TurboFan work in progress*

d8 --experimental-wasm-instruction-tracing --
print-code --liftoff-only test.js

*discussed in issues

000001043AE813E0	0	55	push rbp
000001043AE813E1	1	4889e5	REX.W movq rbp, rsp
000001043AE813E4	4	6a08	push 0x8
000001043AE813E6	6	4881ec08000000	REX.W subq rsp, 0x8
000001043AE813ED	d	488975f0	REX.W movq [rbp-0x10], rsi
000001043AE813F1	11	488b4e23	REX.W movq rcx, [rsi+0x23]
000001043AE813F5	15	483b21	REX.W cmpq rsp, [rcx]
000001043AE813F8	18	0f864a000000	jna 000001043AE81448 <+0x68>
000001043AE813FE	1e	53	push rbx
000001043AE813FF	1f	bb11000000	movl rbx, 0000000000000011
000001043AE81404	24	64	'Unimplemented Instruction'
000001043AE81405	25	67	'Unimplemented Instruction'
000001043AE81406	26	90	nop
000001043AE81407	27	90	nop
000001043AE81408	28	90	nop
000001043AE81409	29	5b	pop rbx
000001043AE8140A	2a	53	push rbx
000001043AE8140B	2b	bb12000000	movl rbx, 0000000000000012
000001043AE81410	30	64	'Unimplemented Instruction'
000001043AE81411	31	67	'Unimplemented Instruction'
000001043AE81412	32	90	nop
000001043AE81413	33	90	nop
000001043AE81414	34	90	nop
000001043AE81415	35	5b	pop rbx
000001043AE81416	36	8d0c02	leal rcx, [rdx+rax*1]
000001043AE81419	39	8bc1	movl rax, rcx
000001043AE8141B	3b	48837df808	REX.W cmpq [rbp-0x8], 0x8
000001043AE81420	40	7421	jz 000001043AE81443 <+0x63>
000001043AE81422	42	b924000000	movl rcx, 0000000000000024
000001043AE81427	47	4989e2	REX.W movq r10, rsp
000001043AE8142A	4a	4883ec28	REX.W subq rsp, 0x28
000001043AE8142E	4e	4883e4f0	REX.W andq rsp, 0xf0
000001043AE81432	52	4c89542420	REX.W movq [rsp+0x20], r10
000001043AE81437	57	48b840431335f77f0000	REX.W movq rax, 00007FF735134340
000001043AE81441	61	ffdf	call rax
000001043AE81443	63	488be5	REX.W movq rsp, rbp
000001043AE81446	66	5d	pop rbp
000001043AE81447	67	c3	retl
000001043AE81448	68	50	push rax
000001043AE81449	69	52	push rdx
000001043AE8144A	6a	56	push rsi
000001043AE8144B	6b	e830feffff	call 000001043AE81280
000001043AE81450	70	5e	pop rsi
000001043AE81451	71	5a	pop rdx
000001043AE81452	72	58	pop rax
000001043AE81453	73	eba9	jmp 000001043AE813FE <+0x1e>
000001043AE81455	75	0f1f00	nop

Current Issues

- Clang
 - There are certain optimizations after the translation to LLVM IR that cause the trace instructions to be moved to the top of the function. The LLVM IR generates correctly, but as you can see in the ASM this is not the case.
- V8
 - In Liftoff, trace instruction is emitted where the offset is.
 - In TurboFan, inside of a function if there are more than one trace instruction, all but the first get optimized away.
- Emscripten
 - The compilation works fine (just clang), but the optimization step with binaryen sometimes causes problems. When the WASM gets reordered, the instTrace offsets don't move with it.

Summary

- Instruction Tracing
 - Custom section defines where trace instructions should be emitted
 - Decision made at the engine implementation
 - Compiler intrinsic to generate custom section
- Vote for Phase 2