

# Introduction

## WebAssembly

What makes WebAssembly different from Assembly?

- Foundational concept is the Module, not the MMU
- I/O goes through imports and exports, not syscalls
  - Declared
  - Typed

## Virtualization 🤖

## Shared-nothing linking


Isolation without virtual address space boundaries

Calls as "IPC"

## Virtualization 🤖

## The hardest problem in computers 😊

How do I copy data from one computer to another? 🤖



I LIKE HOW WE'VE HAD THE INTERNET FOR DECADES,  
YET "SENDING FILES" IS SOMETHING EARLY  
ADOPTERS ARE STILL FIGURING OUT HOW TO DO.


Source: <https://xkcd.com/949/>

In a similar way, filesystems and socket APIs involve a lot of knowledge of the host.

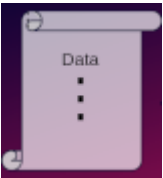

## What's in a File?

Data

Metadata



## Compute and Metacompute

Compute	Metacompute
	
operate on data	configure resources
algorithms and data structures	organize and navigate storage
business logic	manage access between different users
arithmetic	interface with humans

## POSIX is dynamically typed

Quiz! 🤖



What's the difference between this:

```
$ foo < content.txt
```

and this:

```
$ cat content.txt | foo
```

Can you `lseek` on stdin in POSIX?

- If it's a normal file redirected with `< :` 
- If it's a pipe: 

WebAssembly has a static typesystem. Let's use it!

## Virtualization

# I/O Types

## Wasm-native I/O

# I/O Types

## Virtualization

# I/O Types

What is an I/O Type?

- An abstraction for a common I/O pattern
- Expose data, hide metadata

Strategy: Define a few simple I/O types now, add more over time.

# I/O Streams

Overview:

- Reliable in-order delivery of bytes
  - or in the future, of arbitrary type
- Examples: sockets, pipes, character devices, files\*
- Prototype: <https://crates.io/crates/io-streams>, <https://crates.io/crates/nameless>

### Operations:

- `read` (input)
- `write` (output)
- `skip` (input) - like `read` but discards data
- `write_zeros` (output) - like `write` but just writes zeros
- `flush` (output) - flush buffers, report errors
- `media_type` - return a `string` with the IANA Media Type (eg. "image/jpeg")
- `pseudonym` - return a handle to an identifier
- `write_pseudonym` - (output) write the name of a pseudonym to the output
- `forward` - `read` from an input and `write` the data to another ➡

### Constructors:

- `pipe`
- `null` - (output) bit bucket
- `from_array` - form a stream from an I/O array at a given offset

## I/O Arrays

### Overview:

- Dynamic array of bytes
  - or in the future, of arbitrary type
- Examples: disk partitions, memory buffers, device memory, files\*
- Prototype: <https://crates.io/crates/io-arrays>

### Operations:

- `read_at`
- `write_at`
- `len` - current length
- `set_len` - set current length, truncating or appending zeros
- `advise` - optimization hints
- `flush` - flush buffers, report errors
- `media_type` - return a `string` with the IANA Media Type (eg. "image/jpeg")
- `pseudonym` - return a handle to an identifier
- `copy` - `read_at` from one I/O array and `write_at` the data to another ➡

### Constructors:

- `anonymous`
- `from_stream` - read all the data from an input stream and write it into an anonymous I/O array

## Typed Main

Typed main functions will be able to take I/O Types as arguments.

Programs which operate this way are independent of:

- Network stacks
- Filesystems
- Case sensitivity, Unicode normalization, path separators, the behavior of `..`, how symlinks work, non-Unicode filenames, filename length

limitations, trailing whitespace, ...

## Errors

Report partial success along with failure.

- Wasm has multiple return values!

Error types:

- Applications usually don't care why `read` or `write` failed
- Detailed errors expose host information
- Options:
  - Success/Failure
  - Opaque handle

No `EINTR`

- Wasm has no signals! 😎
- No need to wrap `read` and `write` in retry loops

Don't be a file or a socket API:

- Errors and `EOF` are sticky!
- No `ETIMEDOUT`
- Atomicity? Seek higher-level APIs that preserve intent

## Relationship to filesystem and sockets

The wasi-filessystem and (anticipated) wasi-sockets proposals will continue:

- Now with less pressure to provide perfect portability
- Can now more easily expand in scope

Some use cases will always need the extra functionality of wasi-filessystem and wasi-sockets.

Some host environments won't support that functionality.

WASI can accommodate both.

## Other WASI APIS

wasi-clocks and wasi-random - convert to handles

- Add a Clock I/O type?

wasi-poll - add support for polling I/O Types

## In Depth



## Are files I/O Streams or I/O Arrays?

Both! (like POSIX)

With separate handles! (unlike POSIX)

What about `lseek`?

- POSIX `lseek` bridges between array and stream
- Emulate `lseek` in libc

## What about directories?

Directories have wildly differing semantics between platforms.

Alternative: stream of lazy streams

# Write-Once Run Anywhere?

What's different between I/O Types and other common WORA approaches?

- Start with compute, not whole applications
- Push metacompute to the boundaries
- At the boundaries, maybe we still have filenames, or even GUIs

## Compatibility with existing code

A libc mode that uses I/O streams and I/O arrays will be able to support:

- `read`, `write`, `pread`, `pwrite`
  - and everything build around them
- all non-I/O things
- `lseek`
- Emulation for `open`, `mmap`, `socket`, etc.

Continue to support the libc mode that uses `wasi-filessystem` and `wasi-sockets`.

Let developers chose!

## Zero copy

- Ways to do zero-copy
  - `forward`, `skip`, `write_zeros`, `copy`
  - `io-array`'s `read_at` skips intervening data
  - create an I/O Array and pass around a handle
- `read` is not zero-copy
  - shared pages require protocols, synchronization
    - not virtualizable
- What can we do instead?
  - Avoid bringing data into an instance that doesn't need it
  - Use direct calls instead of shared-memory protocols where appropriate

## Async

WebAssembly is starting a subgroup to talk about stack switching:



<https://github.com/WebAssembly/meetings/blob/master/main/2021/CG-03-16.md#agenda-items>

[https://docs.google.com/presentation/d/1tel1N282B7Dog314G22Emmaur2XCVtqYwwq5c\\_AK-BE/edit#slide=id.p](https://docs.google.com/presentation/d/1tel1N282B7Dog314G22Emmaur2XCVtqYwwq5c_AK-BE/edit#slide=id.p)

Like the rest of WASI, I/O Types are independent of sync/async.

## Data parallelism

That's a separate talk!

## Things to come

- Vectored I/O - Maybe interface types can enable this automatically
- Multiplexed connections (QUIC, etc.)
- I/O Windows
- Interactive streams
- Text streams
- Terminal I/O
- A way to do `fsync` / `fdatsync` which preserves intent